



Project | Deep Learning: Image Classification with CNN

Collaborators:


Ginosca Alejandro
Natanael Santiago



Project Overview

- Building a Convolutional Neural Network (CNN)
 - Classify images
-
- Animals-10 Dataset
 - ~ 28K images
 - Different sizes
 - 10 classes:

Dataset preparation





Dataset preparation

- Data exploration
 - Smallest dimensions (60, 57),
 - Largest dimensions (6720, 4480)
 - Approximate average dimension (326, 326)
- Preprocessing
 - Resizing (128 x 128) and normalizing
 - Training set: 80% (20,947)
 - Validation Set: 20% (5,232)

Images after pre-processing

Class: spider



Class: cat



Class: butterfly



Class: chicken



Class: sheep



Class: squirrel



Class: dog



Class: cow



Class: horse



Class: elephant






CNN

- Architecture:
 - Convolutional and MaxPooling layers were stacked to extract features from the input images.
 - Fully connected dense layers were used for classification.
 - A final softmax layer outputs predictions for the 10 classes.
- Key Design Choices:
 - Simplicity in design for computational efficiency.
 - Moderate depth to avoid overfitting.

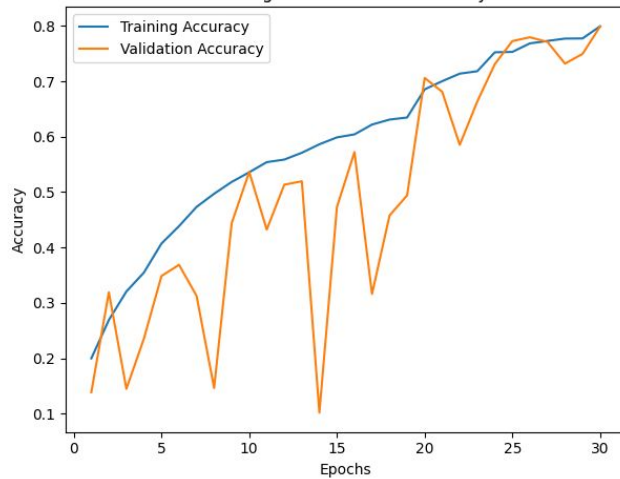
Learning rate	Batch Size	Epochs	Callbacks
0.001	32	30	EarlyStopping ReduceLROnPlateau

CNN

Training and Validation Loss

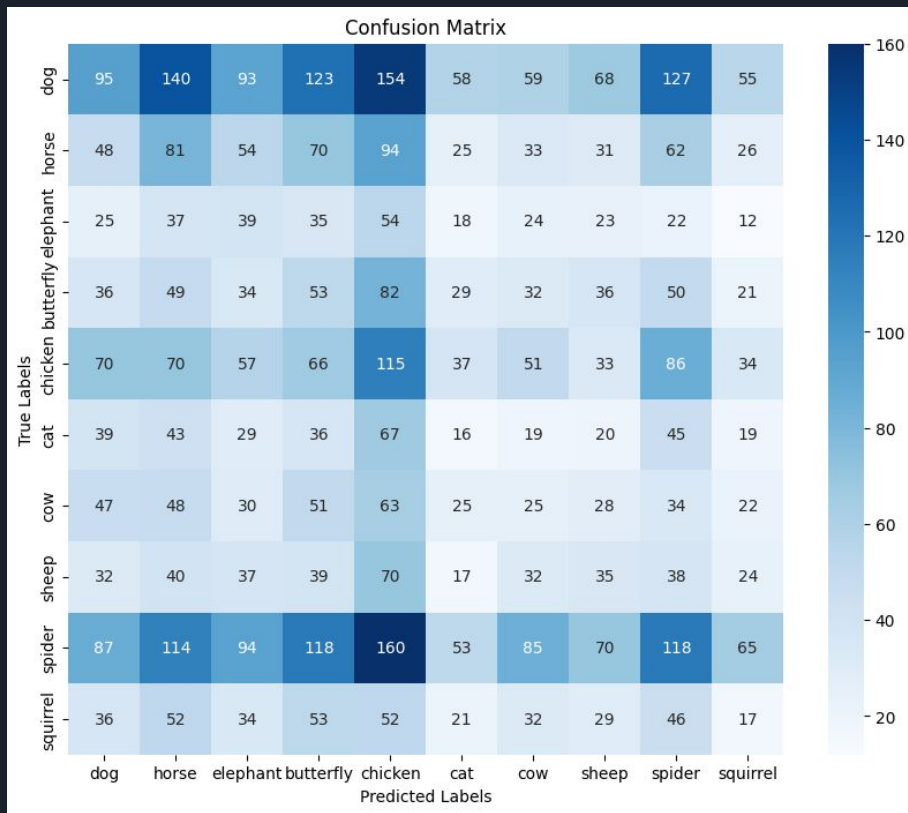


Training and Validation Accuracy



Metric	Custom Model
Validation Loss	1.0980
Validation Accuracy	72.32%
Validation Precision	78.12%
Validation Recall	67.62%

CNN



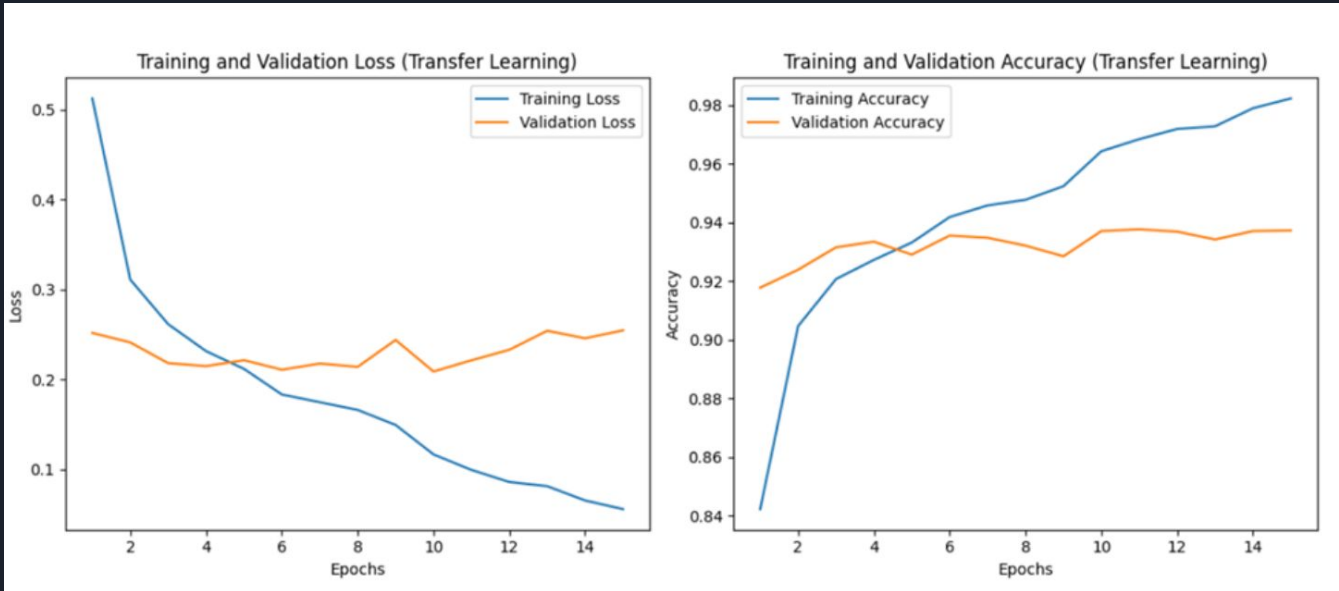


Transfer Learning

- MobileNetV2 pretrained on ImageNet.

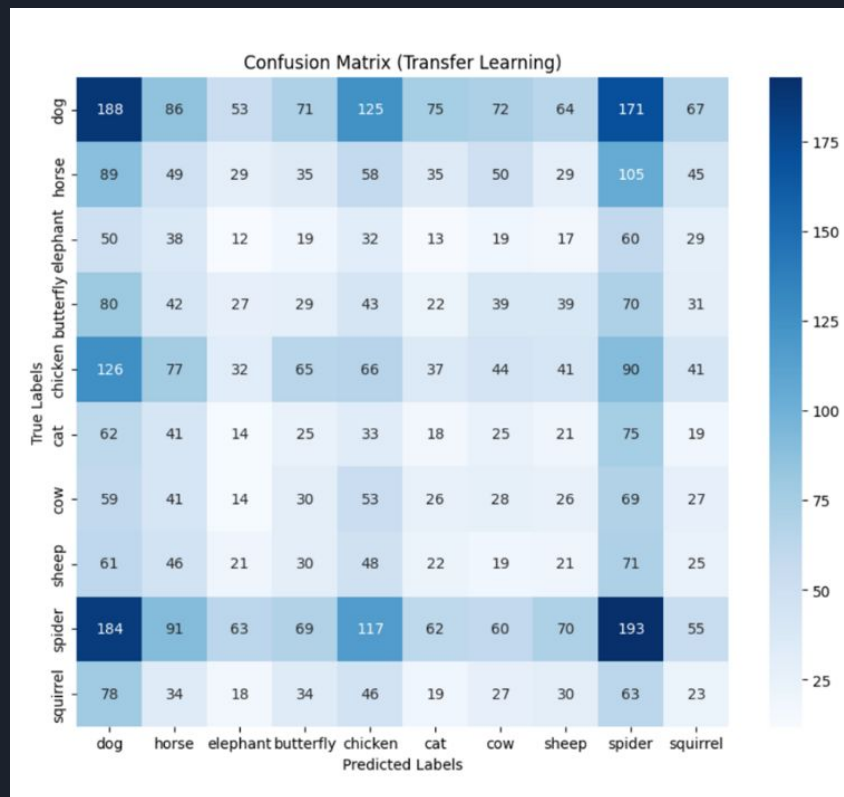
```
# Add custom classification layers
x = base_model.output
x = GlobalAveragePooling2D()(x) # Global Average Pooling layer
x = Dense(256, activation='relu')(x) # Fully connected layer
x = Dropout(0.5)(x) # Dropout for regularization
x = Dense(128, activation='relu')(x) # Fully connected layer
x = Dropout(0.3)(x) # Dropout for regularization
output = Dense(10, activation='softmax')(x) # Output layer for 10 classes
```

Transfer Learning



Validation Loss	0.2086
Validation Accuracy	93.71%
Validation Precision	95.29%
Validation Recall	92.78%

Transfer Learning







Comparison

Metric	Custom Model	Transfer Learning Model
Validation Loss	1.0980	0.2086
Validation Accuracy	72.32%	93.71%
Validation Precision	78.12%	95.29%
Validation Recall	67.62%	92.78%

Comparison (CNN vs Transfer Learning)





Insights and Learnings

- Impact of Transfer Learning:
 - Leveraging pretrained models significantly improved performance compared to building a model from scratch.
- Importance of Preprocessing:
 - Data augmentation and normalization played a crucial role in achieving generalization.
- Challenges with Specific Classes:
 - Categories like "dog" and "butterfly" require more refined feature extraction due to visual similarities.
- Role of Callbacks:
 - EarlyStopping and ReduceLROnPlateau contributed to efficient training and prevented overfitting.



Next Steps

1. Fine-Tuning:
 - Unfreeze additional layers in MobileNetV2 for task-specific learning.
2. Dataset Enhancements:
 - Collect more diverse examples for underperforming classes.
3. Exploring Alternate Architectures:
 - Experiment with other lightweight architectures like EfficientNet or ResNet.
4. Deploying the Model:
 - Integrate the model into an application for real-time image classification.



THANK YOU!

